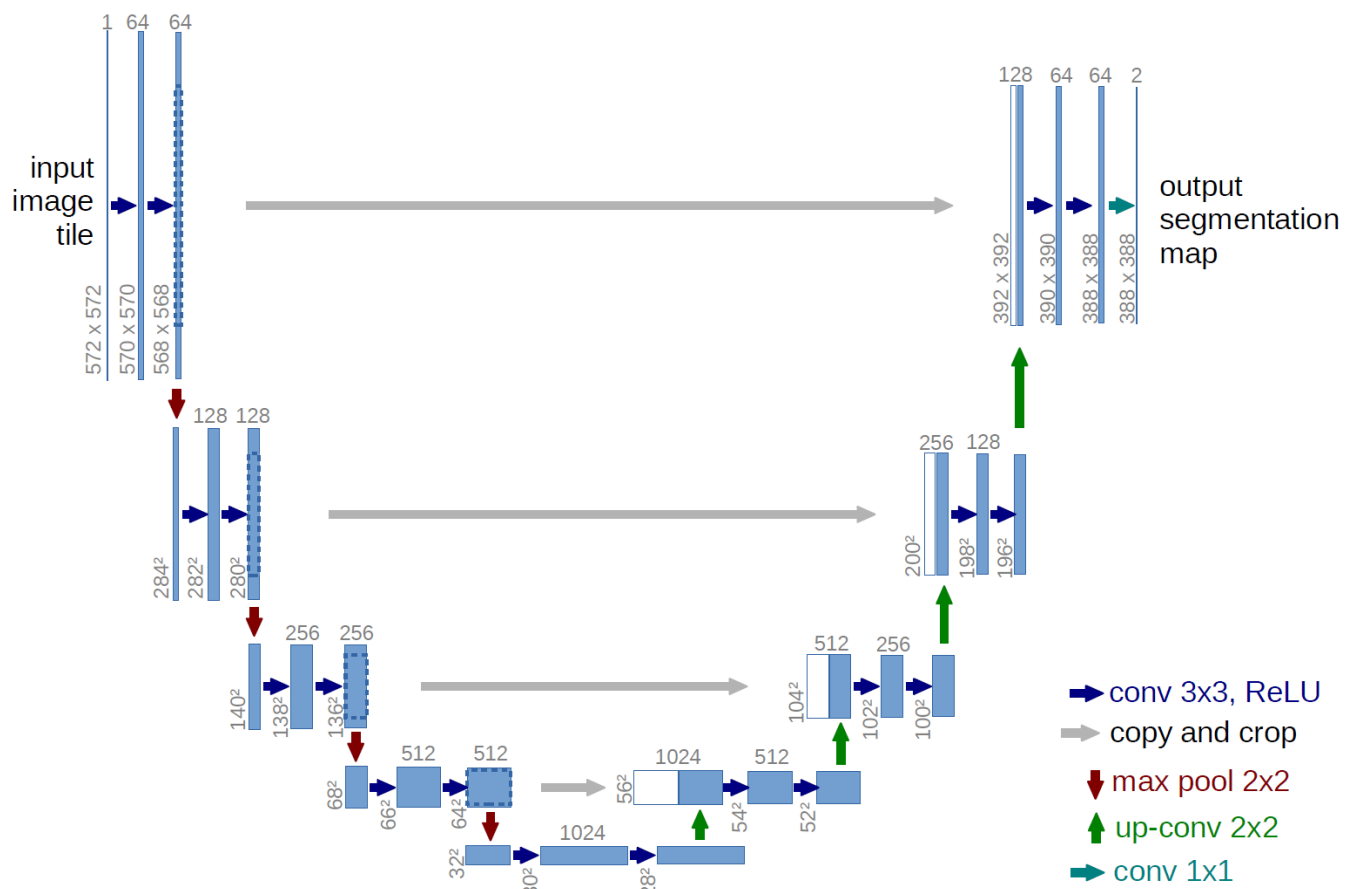


U-Net for cars segmentation

In this work I use U-Net to solve the problem of semantic segmentation: 0 — not a car, 1 — car.



In []:

```
# !unzip /data/rvgorb/hw11/data.zip
```

TensorBoard

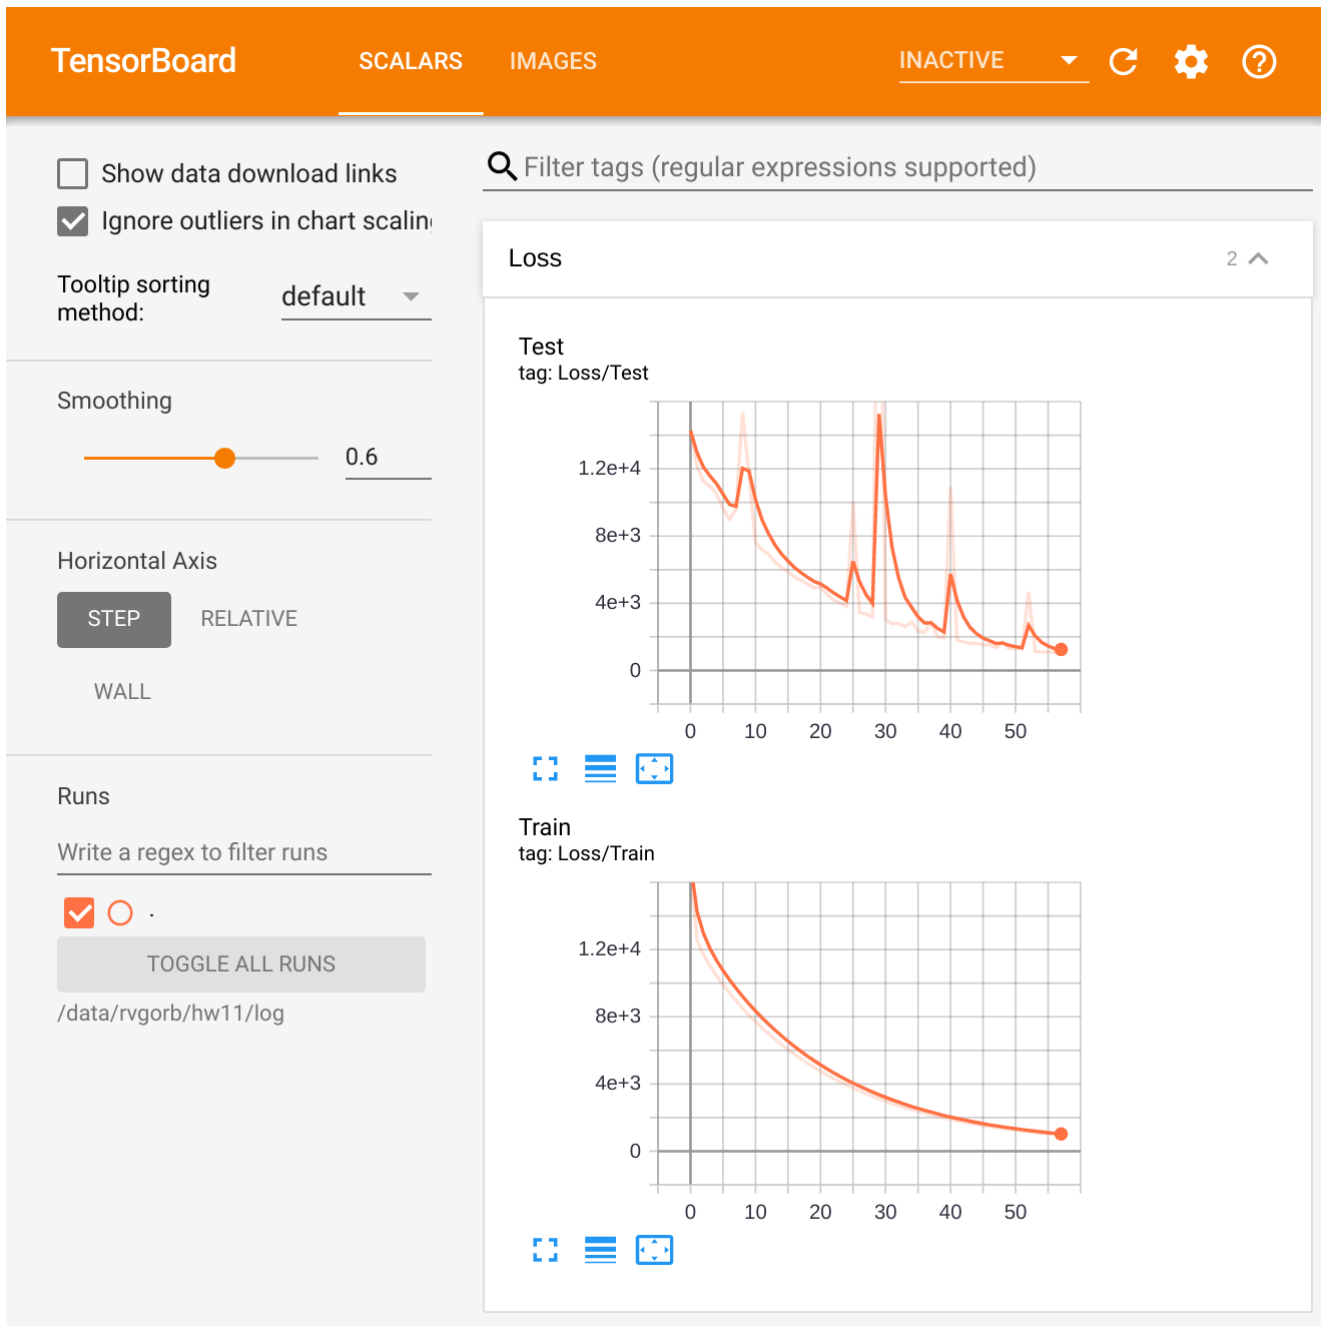
In [168]:

```
!mkdir /data/rvgorb/hw11/log  
%load_ext tensorboard  
%tensorboard --logdir /data/rvgorb/hw11/log --port 7773
```

```
mkdir: cannot create directory '/data/rvgorb/hw11/log': File exists
The tensorboard extension is already loaded. To reload it, use:
  %reload_ext tensorboard
```

```
Reusing TensorBoard on port 7773 (pid 9891), started 0:00:08 ago. (Use '!kill 9891' to kill it.)
```





Training model

In [85]:

```
!python3 train.py --n_epochs 300 --batch_size 20 --learning_rate 0.00001
```

```
Current epoch: 0
100%|██████████████████████████████████████| 229/229 [02:12<00:00]
0, 1.72it/s]
Epoch loss 17139.0282716976
Make test
100%|██████████████████████████████████████| 26/26 [00:07<00:00]
0, 3.37it/s]
Test loss 14279.07421875
Current epoch: 1
100%|██████████████████████████████████████| 229/229 [02:09<00:00]
0, 1.76it/s]
Epoch loss 12553.476886599345
Make test
100%|██████████████████████████████████████| 26/26 [00:08<00:00]
0, 3.13it/s]
Test loss 12178.2054933563
Current epoch: 2
100%|██████████████████████████████████████| 229/229 [02:10<00:00]
0, 1.76it/s]
Epoch loss 11687.493869405022
Make test
100%|██████████████████████████████████████| 26/26 [00:07<00:00]
0, 3.42it/s]
Test loss 11254.069451279527
Current epoch: 3
100%|██████████████████████████████████████| 229/229 [08:10<00:00]
0, 2.14s/it]
Epoch loss 11027.155598389738
Make test
100%|██████████████████████████████████████| 26/26 [00:08<00:00]
0, 3.05it/s]
Test loss 10961.254890501968
Current epoch: 4
100%|██████████████████████████████████████| 229/229 [06:26<00:00]
0, 1.69s/it]
Epoch loss 10443.98858146834
Make test
100%|██████████████████████████████████████| 26/26 [00:14<00:00]
0, 1.80it/s]
Test loss 10518.217658095473
Current epoch: 5
100%|██████████████████████████████████████| 229/229 [02:10<00:00]
0, 1.75it/s]
Epoch loss 9905.598956058951
Make test
100%|██████████████████████████████████████| 26/26 [00:09<00:00]
0, 2.82it/s]
Test loss 9637.588659571851
Current epoch: 6
100%|██████████████████████████████████████| 229/229 [02:11<00:00]
0, 1.75it/s]
Epoch loss 9413.685978438865
Make test
100%|██████████████████████████████████████| 26/26 [00:11<00:00]
0, 2.33it/s]
Test loss 8997.551350270669
Current epoch: 7
100%|██████████████████████████████████████| 229/229 [02:11<00:00]
0, 1.74it/s]
Epoch loss 8954.252746315502
Make test
```

```
100%|██████████| 26/26 [00:07<00:00]
0, 3.33it/s]
Test loss 9596.471702755905
Current epoch: 8
100%|██████████| 229/229 [02:09<00:00]
0, 1.76it/s]
Epoch loss 8509.588659934498
Make test
100%|██████████| 26/26 [00:10<00:00]
0, 2.45it/s]
Test loss 15376.019454355315
Current epoch: 9
100%|██████████| 229/229 [02:09<00:00]
0, 1.76it/s]
Epoch loss 8088.47044555131
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.46it/s]
Test loss 11631.133919783464
Current epoch: 10
100%|██████████| 229/229 [04:13<00:00]
0, 1.11s/it]
Epoch loss 7715.233614219433
Make test
100%|██████████| 26/26 [00:08<00:00]
0, 2.98it/s]
Test loss 7628.0539800688975
Current epoch: 11
100%|██████████| 229/229 [02:10<00:00]
0, 1.76it/s]
Epoch loss 7347.467381959607
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.57it/s]
Test loss 7177.38019039124
Current epoch: 12
100%|██████████| 229/229 [02:09<00:00]
0, 1.76it/s]
Epoch loss 7004.710101664848
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.68it/s]
Test loss 6926.01711675689
Current epoch: 13
100%|██████████| 229/229 [02:09<00:00]
0, 1.76it/s]
Epoch loss 6670.103430335698
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.56it/s]
Test loss 6437.695220226378
Current epoch: 14
100%|██████████| 229/229 [02:09<00:00]
0, 1.76it/s]
Epoch loss 6349.797449849891
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.66it/s]
Test loss 6131.300796628937
Current epoch: 15
100%|██████████| 229/229 [02:09<00:00]
```

[illegible]


```
Make test  
100%|██████████| 26/26 [00:07<00:00]  
0, 3.53it/s]  
Test loss 3016.5042984128936  
Current epoch: 31  
100%|██████████| 229/229 [02:09<00:00]  
0, 1.77it/s]  
Epoch loss 2840.8347792712884  
Make test  
100%|██████████| 26/26 [00:07<00:00]  
0, 3.58it/s]  
Test loss 2787.6638817974904  
Current epoch: 32  
100%|██████████| 229/229 [02:09<00:00]  
0, 1.77it/s]  
Epoch loss 2710.889703875546  
Make test  
100%|██████████| 26/26 [00:07<00:00]  
0, 3.40it/s]  
Test loss 2815.934708722933  
Current epoch: 33  
100%|██████████| 229/229 [02:09<00:00]  
0, 1.76it/s]  
Epoch loss 2589.024857566867  
Make test  
100%|██████████| 26/26 [00:07<00:00]  
0, 3.51it/s]  
Test loss 2608.8815821850394  
Current epoch: 34  
100%|██████████| 229/229 [02:09<00:00]  
0, 1.76it/s]  
Epoch loss 2474.8730903725436  
Make test  
100%|██████████| 26/26 [00:07<00:00]  
0, 3.47it/s]  
Test loss 2878.3261488065946  
Current epoch: 35  
100%|██████████| 229/229 [02:09<00:00]  
0, 1.76it/s]  
Epoch loss 2368.1971197802945  
Make test  
100%|██████████| 26/26 [00:07<00:00]  
0, 3.47it/s]  
Test loss 2331.8658995140254  
Current epoch: 36  
100%|██████████| 229/229 [02:10<00:00]  
0, 1.76it/s]  
Epoch loss 2262.934787800218  
Make test  
100%|██████████| 26/26 [00:07<00:00]  
0, 3.31it/s]  
Test loss 2267.983013964075  
Current epoch: 37  
100%|██████████| 229/229 [02:10<00:00]  
0, 1.76it/s]  
Epoch loss 2165.288143934225  
Make test  
100%|██████████| 26/26 [00:07<00:00]  
0, 3.26it/s]  
Test loss 2863.8117733452264  
Current epoch: 38
```

```
100%|██████████| 229/229 [02:10<00:00]
0, 1.76it/s]
Epoch loss 2070.647650279749
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.58it/s]
Test loss 2004.1658464566929
Current epoch: 39
100%|██████████| 229/229 [02:10<00:00]
0, 1.76it/s]
Epoch loss 1980.331762588701
Make test
100%|██████████| 26/26 [00:09<00:00]
0, 2.89it/s]
Test loss 1960.6812215489665
Current epoch: 40
100%|██████████| 229/229 [02:10<00:00]
0, 1.76it/s]
Epoch loss 1897.90754639738
Make test
100%|██████████| 26/26 [00:08<00:00]
0, 3.10it/s]
Test loss 10936.80144254429
Current epoch: 41
100%|██████████| 229/229 [02:10<00:00]
0, 1.75it/s]
Epoch loss 1815.8012810453056
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.45it/s]
Test loss 1809.693588136688
Current epoch: 42
100%|██████████| 229/229 [02:10<00:00]
0, 1.76it/s]
Epoch loss 1737.1326052469979
Make test
100%|██████████| 26/26 [00:08<00:00]
0, 3.21it/s]
Test loss 1698.5323688176672
Current epoch: 43
100%|██████████| 229/229 [02:10<00:00]
0, 1.75it/s]
Epoch loss 1667.940561885917
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.40it/s]
Test loss 1622.9848440575788
Current epoch: 44
100%|██████████| 229/229 [02:09<00:00]
0, 1.76it/s]
Epoch loss 1598.4927521151747
Make test
100%|██████████| 26/26 [00:07<00:00]
0, 3.35it/s]
Test loss 1625.4190644992618
Current epoch: 45
100%|██████████| 229/229 [02:10<00:00]
0, 1.76it/s]
Epoch loss 1530.7861217248908
Make test
100%|██████████| 26/26 [00:07<00:00]
```

[illegible]

In [86]:

```
# Imports
%load_ext autoreload
%autoreload 2

import matplotlib.pyplot as plt
import seaborn as sns
import torch
import torch.utils.data as dt
from tqdm.auto import tqdm

from carvana_dataset import CarvanaDataset
from conf import *
from model import SegmenterModel
from stats import *
```

The autoreload extension is already loaded. To reload it, use:
%reload_ext autoreload

In [110]:

```
# Datasets
train_data = CarvanaDataset('/data/rvgorb/hw11/train/',
                             '/data/rvgorb/hw11/train_masks/')
test_data = CarvanaDataset('/data/rvgorb/hw11/test/',
                             '/data/rvgorb/hw11/test_masks/',
                             is_train=False)
```

Loading trained model:

In [111]:

```
unet = SegmenterModel(3, 1).cuda()
unet.load_state_dict(
    torch.load("/data/rvgorb/hw11/unet_dump_recent", map_location="cuda:6"))
unet.train(False)
```

In [162]:

```
def plot_predictions_on_random_image(fitted_model: SegmenterModel,
                                     test_data: CarvanaDataset,
                                     seed: int = None):
    if seed is None:
        rvgen = np.random
    else:
        rvgen = np.RandomState(seed)

    n_img = len(test_data.files)
    img_idx = rvgen.randint(0, n_img)
    img = test_data[img_idx][0]
    true_mask = test_data[img_idx][1].squeeze(0)

    pred_logits = fitted_model.forward(img.unsqueeze(0).cuda())
    pred_logits = pred_logits\
        .squeeze(0).squeeze(0)\
        .detach().cpu().numpy()
    pred_mask = pred_logits > 0.01

    images_to_plot = [true_mask, pred_mask, pred_logits]
    ax_titles = ["Ground truth mask",
                 "Predicted mask (threshold = 0.01)",
                 "Predicted logits"]


    fig, axes = plt.subplots(1, 3, figsize=(20, 10))
    for i in range(len(axes)):
        ax = axes[i]
        to_plot = images_to_plot[i]
        ax.imshow(to_plot, cmap="inferno")
        ax.set_title(ax_titles[i], fontsize=15, fontweight="bold")
        ax.set_axis_off()
```

Predicted and ground truth masks comparison:

In [163]:

```
for i in range(5):  
    plot_predictions_on_random_image(unet, test_data)
```



Ground truth mask



Predicted mask (threshold — 0.01)



Predicted logits




Ground truth mask



Predicted mask (threshold — 0.01)



Predicted logits



Ground truth mask



Predicted mask (threshold — 0.01)




Predicted logits




Ground truth mask




Predicted mask (threshold — 0.01)



Predicted logits





Quality estimation

IoU — *Intersection over Union* — is an evaluation metric used to measure the accuracy of an object detector on a particular dataset.

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$

Examining this equation you can see that Intersection over Union is simply a **ratio**.

In the numerator we compute the **area of overlap** between the predicted bounding box and the ground-truth bounding box.

The denominator is the **area of union**, or more simply, the area encompassed by both the predicted bounding box and the ground-truth bounding box

In [154]:

```
def iou_loss(prediction, ground_truth):
    """ Считает IoU по всем элементам батча """
    n_images = len(prediction)
    intersection, union = 0, 0
    for i in range(n_images):
        intersection += np.logical_and(
            prediction[i] > 0, ground_truth[i] > 0).astype(np.float32).sum()
        union += np.logical_or(prediction[i] > 0,
                                ground_truth[i] > 0).astype(np.float32).sum()
    return float(intersection) / union
```

In [159]:


```
# Quality on test
dl_test = dt.DataLoader(test_data, shuffle=False, num_workers=8, batch_size=1)

test_loss = []
for i, (input_batch, target_mask_batch) in enumerate(tqdm(dl_test)):
    with torch.no_grad():
        output_logits_batch = unet(input_batch.to("cuda")).cpu()
        output_mask_batch = output_logits_batch > 0.03
        loss = iou_loss(output_mask_batch.cpu().numpy(),
                        target_mask_batch.cpu().numpy() > 0)
        test_loss.append(loss)
test_loss = np.hstack(test_loss)

plt.figure(figsize=(16, 9))
sns.distplot(test_loss)
_ = plt.title("Distribution of IoU on test")

print(f"Mean value of of IoU on test: {np.mean(test_loss)}")
```

Mean value of of IoU on test: 0.9728439083072745




Quality on radom photo from web:

In [94]:

```
# !wget https://i.pinimg.com/originals/2e/49/11/2e49114884cf3d543454a687fb55b3e7.jpg
# !mkdir example
# !mv 2e49114884cf3d543454a687fb55b3e7.jpg example/
# !mv 2e49114884cf3d543454a687fb55b3e7.gif example_mask/
```

In [164]:

```
example_data = CarvanaDataset('./example/', './example_mask/', is_train=False)
plot_predictions_on_random_image(unet, example_data)
```



In [166]:

```
# IoU computation
dl_example = dt.DataLoader(example_data, shuffle=False,
                           num_workers=8,
                           batch_size=1)

test_loss = []
for i, (input_batch, target_mask_batch) in enumerate(tqdm(dl_example)):
    with torch.no_grad():
        output_logits_batch = unet(input_batch.to("cuda:0")).cpu()
        output_mask_batch = output_logits_batch > 0.
        loss = iou_loss(output_mask_batch.cpu().numpy(), target_mask_batch.cpu().numpy() > 0)
    test_loss.append(loss)

test_loss = np.hstack(test_loss)
print(f"IoU: {np.mean(test_loss)}")
```

IoU: 0.9357152344268828

Conclusions:

- Achieved 98% IoU on test, which is not bad. On random image from web got a little bit worse, because the image contains unseen patterns like reflection from the floor which bothers the model to detect wheels of the car.